

```
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
```

```
!kaggle datasets download -d salader/dogs-vs-cats
```

```
Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod
Downloading dogs-vs-cats.zip to /content
 98% 1.05G/1.06G [00:08<00:00, 280MB/s]
100% 1.06G/1.06G [00:08<00:00, 139MB/s]
```



```
import zipfile
zip_ref = zipfile.ZipFile('/content/dogs-vs-cats.zip', 'r')
zip_ref.extractall('/content')
zip_ref.close()
```

```
import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, BatchNormalization, Dropout
```

```
#generator : Divides the data into the batches at a time
```

```
train_ds = keras.utils.image_dataset_from_directory(
    directory = '/content/train',
    labels = 'inferred',
    label_mode = 'int',
    batch_size = 32,
    image_size =(256, 256)
)
```

```
validation_ds = keras.utils.image_dataset_from_directory(
    directory = '/content/test',
    labels = 'inferred',
    label_mode = 'int',
    batch_size = 32,
    image_size =(256, 256)
)
```

```
Found 20000 files belonging to 2 classes.
Found 5000 files belonging to 2 classes.
```

```
# Normalize
```

```
def process(image, label):
    image = tf.cast(image/255. ,tf.float32)
    return image, label
```

```
train_ds = train_ds.map(process)
validation_ds = validation_ds.map(process)
```

```
# CNN Model
model = Sequential()

model.add(Conv2D(32, kernel_size=(3,3), padding= 'valid', activation='relu', input_shape=(256,256,3)))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(64, kernel_size=(3,3), padding= 'valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Conv2D(128, kernel_size=(3,3), padding= 'valid', activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2), strides=2, padding='valid'))

model.add(Flatten())

model.add(Dense(128, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.1))
model.add(Dense(1, activation='sigmoid'))

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 254, 254, 32)	896
batch_normalization (Batch Normalization)	(None, 254, 254, 32)	128
max_pooling2d_3 (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_4 (Conv2D)	(None, 125, 125, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 125, 125, 64)	256
max_pooling2d_4 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_5 (Conv2D)	(None, 60, 60, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 60, 60, 128)	512
max_pooling2d_5 (MaxPooling2D)	(None, 30, 30, 128)	0
flatten_1 (Flatten)	(None, 115200)	0
dense_3 (Dense)	(None, 128)	14745728
dropout (Dropout)	(None, 128)	0
dense_4 (Dense)	(None, 64)	8256
dropout_1 (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 1)	65

=====

Total params: 14848193 (56.64 MB)

Trainable params: 14847745 (56.64 MB)

Non-trainable params: 448 (1.75 KB)

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(train_ds, epochs=10, validation_data=validation_ds)
```

```
Epoch 1/10
625/625 [=====] - 74s 110ms/step - loss: 1.3066 - accuracy: 0.5938 - val_loss:
Epoch 2/10
625/625 [=====] - 65s 103ms/step - loss: 0.5973 - accuracy: 0.6832 - val_loss:
Epoch 3/10
625/625 [=====] - 66s 105ms/step - loss: 0.4944 - accuracy: 0.7563 - val_loss:
Epoch 4/10
625/625 [=====] - 66s 105ms/step - loss: 0.4597 - accuracy: 0.7928 - val_loss:
Epoch 5/10
625/625 [=====] - 66s 105ms/step - loss: 0.3897 - accuracy: 0.8233 - val_loss:
Epoch 6/10
625/625 [=====] - 78s 124ms/step - loss: 0.3317 - accuracy: 0.8550 - val_loss:
Epoch 7/10
625/625 [=====] - 67s 107ms/step - loss: 0.2547 - accuracy: 0.8925 - val_loss:
Epoch 8/10
625/625 [=====] - 65s 104ms/step - loss: 0.1857 - accuracy: 0.9268 - val_loss:
Epoch 9/10
625/625 [=====] - 66s 105ms/step - loss: 0.1400 - accuracy: 0.9483 - val_loss:
Epoch 10/10
625/625 [=====] - 65s 104ms/step - loss: 0.1010 - accuracy: 0.9633 - val_loss:
```

```
# Checking for training and validation accuracy
```

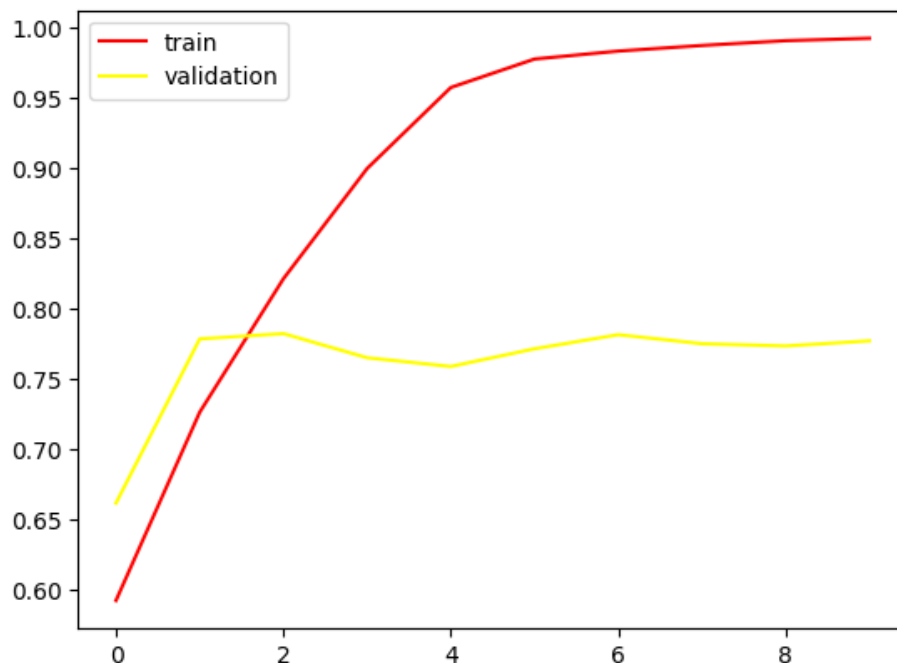
```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['accuracy'], color='red', label='train')
```

```
plt.plot(history.history['val_accuracy'], color='yellow', label='validation')
```

```
plt.legend()
```

```
plt.show()
```



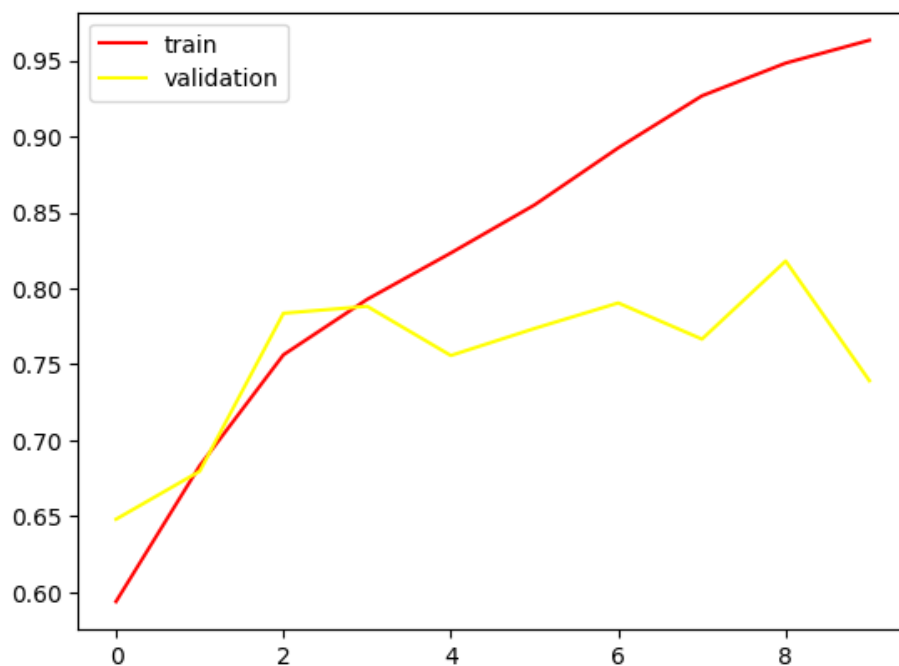
```
# After applying the batchNormalization
```

```
import matplotlib.pyplot as plt
```


```
plt.plot(history.history['accuracy'], color='red', label='train')
```

```
plt.plot(history.history['val_accuracy'], color='yellow', label='validation')
```


```
plt.legend()  
plt.show()
```



```
plt.plot(history.history['loss'], color='red', label='train')  
plt.plot(history.history['val_loss'], color='yellow', label='validation')  
plt.legend()  
plt.show()
```



```
plt.plot(history.history['loss'], color='red', label='train')  
plt.plot(history.history['val_loss'], color='yellow', label='validation')  
plt.legend()  
plt.show()
```



it is a clear indication of overfitting the model